

# CS343 Operating Systems

## Lecture 4

### CPU Scheduling Algorithms




**John Jose**

**Associate Professor**

**Department of Computer Science & Engineering**

**Indian Institute of Technology Guwahati**

# Process State Diagram



# Scheduling Criteria

- ❖ **CPU Utilization** – Percentage time CPU is busy executing process. ↑
- ❖ **Throughput** - Number of processes that are completed per time unit. ↑
- ❖ **Turnaround time** - The interval from the time of submission of a process to the time of completion. ↓
- ❖ **Waiting Time** - Amount of time that a process spends waiting in the ready queue. ↓
- ❖ **Response Time** - Time from the submission of a request until the first response is produced. ↓

# CPU Scheduling Algorithms

## Batch Systems

- ❖ First-come first-served
- ❖ Shortest job first
- ❖ Shortest remaining Time next

## Interactive Systems

- ❖ Round-robin scheduling
- ❖ Priority scheduling
- ❖ Multiple queues
- ❖ Shortest process next
- ❖ Guaranteed scheduling
- ❖ Lottery scheduling
- ❖ Fair-share scheduling

# Round Robin Scheduling

- ❖ Modified version of preemptive FCFS
- ❖ Each process gets a small unit of CPU time (time quantum)
- ❖ FIFO queue is used as input
- ❖ When a process enters/re-enters the ready queue, its PCB is linked the tail of the queue
- ❖ After quantum expires, the process is preempted and added to the tail of the ready queue (Hence, preemptive scheduling algorithm)
- ❖ CPU is allocated to the process at the head of the queue
- ❖ Longer process may have multiple context switch before completion


# Round Robin Scheduling

❖ Consider the following process arriving at  $T_0$ , time quantum of 3 units

❖ P1 burst time: 24

❖ P2 burst time: 3

❖ P3 burst time: 9



❖ Waiting Time

❖ P1:  $(6+3+3) = 12$ , P2: 0, P3:  $(6+3+3) = 12$

❖ Completion Time:

❖ P1: 36, P2: 3, P3: 21

❖ Average Waiting Time:  $(0+12+12)/3 = 8$

❖ Average Completion Time:  $(3+21+36)/3 = 20$

# Round Robin Scheduling

- ❖ RR scheduling is better for short jobs and fair
- ❖ Shorter response time, good for interactive jobs
- ❖ Context-switching time adds up for long jobs
- ❖ Context switching takes additional time and overhead
- ❖ If the chosen quantum is
  - ❖ too large, response time suffers
  - ❖ infinite, performance is the same as FIFO
  - ❖ too small, throughput suffers and percentage overhead grows

# Priority Scheduling

- ❖ Each process has a priority number
- ❖ Highest priority process is scheduled first; if equal priorities, then FCFS
- ❖ Managed with a priority queue with priority value as input
- ❖ When a process enters the ready queue, its PCB is linked onto the priority queue at the appropriate entry
- ❖ CPU is allocated to the process at the head of the queue
- ❖ It can have 2 variants; non-preemptive and preemptive
- ❖ Arrival of a new process with a higher priority can preempt the currently running process.



# Issues with Priority Scheduling

- ❖ Consider a scenario in which there are three processes, a high priority (H), a medium priority (M), and a low priority (L).
- ❖ Process L is running and successfully acquires a resource file.
- ❖ Process H begins; since we are using a preemptive priority scheduler, process L is preempted for process H.
- ❖ Process H tries to acquire L's resource, and blocks (held by L).
- ❖ Process M begins running, and, since it has a higher priority than L, it is the highest priority ready process. It preempts L and runs, thus starving high priority process H.
- ❖ This is known as priority inversion. What can we do?

# Priority Inversion

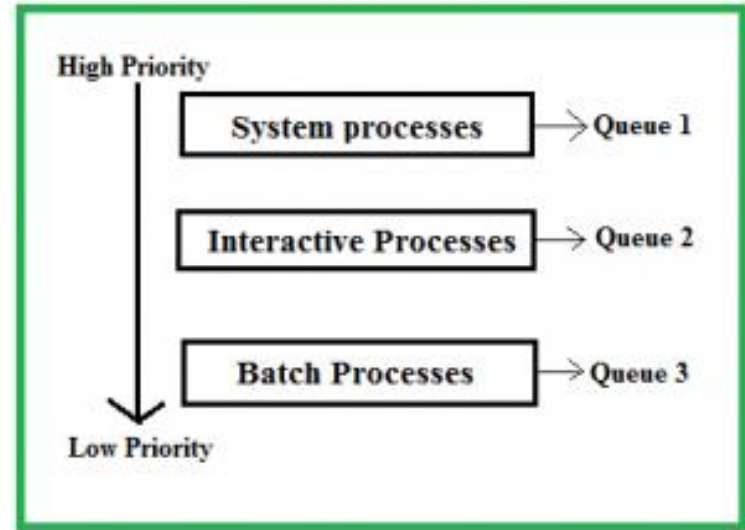
- ❖ Process L should, in fact, be temporarily of higher priority than process M, on behalf of process H.
- ❖ Process H can donate its priority to process L, which, in this case, would make it higher priority than process M.
- ❖ This enables process L to preempt process M and run.
- ❖ When process L is finished, process H becomes unblocked.
- ❖ Process H, now being the highest priority ready process, runs, and process M must wait until it is finished.

# Multilevel Queue

- ❖ Ready queue is partitioned into separate queues:
  - ❖ Foreground, interactive process → RR scheduling
  - ❖ Background, batch process → FCFS scheduling
- ❖ A process is permanently assigned to one queue
- ❖ Each queue has its own scheduling algorithm
- ❖ Can be preemptive

# Multilevel Feedback Queue Scheduling

- ❖ Scheduling must be done between the queues.
  - ❖ Fixed priority scheduling
  - ❖ Serve all from foreground then from background
  - ❖ Possibility of starvation
- ❖ Time slice
  - ❖ Each queue gets a certain amount of CPU time which it can schedule among its processes
    - ❖ i.e.: 80% Vs 20%




# Multilevel Feedback Queue Scheduling

- ❖ **A process can move between the various queues. (Aging)**
- ❖ Multilevel-feedback-queue scheduler defined by the following parameters:
  - ❖ number of queues
  - ❖ scheduling algorithms for each queue
  - ❖ method used to determine when to upgrade a process
  - ❖ method used to determine when to demote a process
  - ❖ method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

## ❖ Three queues:

- ❖ Q0 – time quantum 8 milliseconds, FCFS
- ❖ Q1 – time quantum 16 milliseconds, FCFS
- ❖ Q2 – FCFS



- ❖ A new job enters queue Q0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q1.
- ❖ At Q1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q2.

# Lottery Scheduling

- ❖ Each job some number of lottery tickets are issued
- ❖ On each time slice, randomly pick a winning ticket
- ❖ On average, CPU time is proportional to number of tickets given to each job over time
- ❖ How to assign tickets?
  - ❖ To approximate SRTF, short-running jobs get more, long running jobs get fewer
  - ❖ To avoid starvation, every job gets at least one ticket (everyone makes progress)

# Example: Lottery Scheduling


- ❖ Assume short jobs get 10 tickets, long jobs get 1 ticket

# short jobs / # long jobs	% of CPU each short job gets	% of CPU each long job gets
1/1	91%	9%
0/2	N/A	50%
2/0	50%	N/A
10/1	9.9%	0.99%
1/10	50%	5%



# Conclusion

- ❖ Scheduling: selecting a waiting process from the ready queue and allocating the CPU to it
- ❖ When do the details of the scheduling policy and fairness really matter?
  - ❖ When there aren't enough resources to go around



# Conclusion

- ❖ FCFS scheduling, FIFO Run Until Done:
  - ❖ Simple, but short jobs get stuck behind long ones
- ❖ RR scheduling:
  - ❖ Give each thread a small amount of CPU time when it executes, and cycle between all ready threads
  - ❖ Better for short jobs, but poor when jobs are the same length
- ❖ SJF/SRTF:
  - ❖ Run whatever job has the least amount of computation to do / least amount of remaining computation to do
  - ❖ Optimal (average response time), but unfair; hard to predict the future

# Conclusion

- ❖ Multi-Level Feedback Scheduling:
  - ❖ Multiple queues of different priorities
  - ❖ Automatic promotion/demotion of process priority to approximate SJF/SRTF
- ❖ Lottery Scheduling:
  - ❖ Give each thread a number of tickets (short tasks get more)
  - ❖ Every thread gets tickets to ensure forward progress / fairness
- ❖ Priority Scheduling:
  - ❖ Preemptive or Non-preemptive
  - ❖ Priority Inversion



**johnjose@iitg.ac.in**  
**<http://www.iitg.ac.in/johnjose/>**