

CS343 Operating Systems

Lecture 11

Introduction to Process Deadlocks



John Jose

Associate Professor

Department of Computer Science & Engineering

Indian Institute of Technology Guwahati

Objectives of Deadlock Management Unit

- ❖ To describe deadlocks, which prevent sets of concurrent processes from completing their tasks
- ❖ To discuss different methods for preventing or avoiding deadlocks in a computer system

System Model

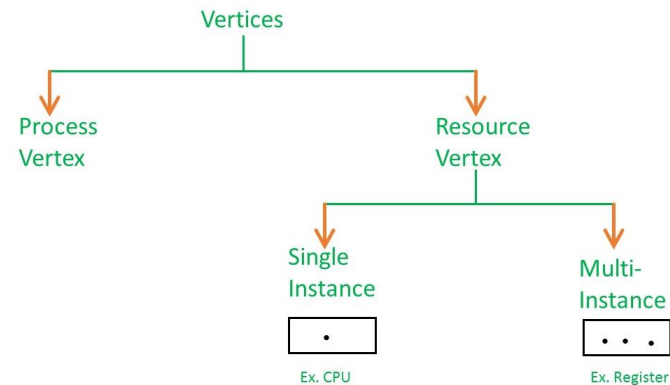
- ❖ System consists of resources
- ❖ Resource types R_1, R_2, \dots, R_m
 - ❖ *CPU cycles, memory space, I/O devices*
- ❖ Each resource type R_i has W_i instances.
- ❖ Each process utilizes a resource as follows:
 - ❖ **request**
 - ❖ **use**
 - ❖ **release**

Deadlock Characterization

- ❖ Deadlock can arise if the following four conditions hold simultaneously.
- ❖ **Mutual exclusion:** Only one process at a time can use a resource
- ❖ **Hold and wait:** A process holding at least one resource is waiting to acquire additional resources held by other processes
- ❖ **No preemption:** A resource can be released only voluntarily by the process holding it, after that process has completed its task
- ❖ **Circular wait:** There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Resource-Allocation Graph

- ❖ A set of vertices V and a set of edges E .
- ❖ V is partitioned into two types:
 - ❖ $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the active processes in the system
 - ❖ $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- ❖ **request edge** – directed edge $P_i \rightarrow R_j$
- ❖ **assignment edge** – directed edge $R_j \rightarrow P_i$



Resource-Allocation Graph


❖ Process





❖ Resource Type with 4 instances





❖ P_i requests an instance of R_j




❖ P_i is holding an instance of R_j



Resource-Allocation Graph



RAG with a deadlock




RAG without a deadlock


Deadlock detection in RAG

- ❖ If graph contains no cycles \Rightarrow no deadlock
- ❖ If graph contains a cycle \Rightarrow
 - ❖ if only one instance per resource type, then deadlock
 - ❖ if several instances per resource type, possibility of deadlock

Deadlock detection in RAG



**A cycle...and
deadlock!**



**Same cycle...but no
deadlock. Why?**

Methods for Handling Deadlocks

- ❖ Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX
- ❖ Ensure that the system will **never** enter a deadlock state:
 - ❖ Deadlock prevention
 - ❖ Deadlock avoidance
- ❖ Allow the system to enter a deadlock state and then recover

Deadlock Prevention

- ❖ **Deadlock prevention** is done by ensuring that at least one of the necessary 4 conditions for deadlock is not met.
- ❖ **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources
- ❖ **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - ❖ Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none allocated to it.
 - ❖ Low resource utilization; starvation possible

Deadlock Prevention

❖ No Preemption –

- ❖ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
 - ❖ Preempted resources are added to the list of resources for which the process is waiting
 - ❖ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting
- ## ❖ Circular Wait –
- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

Deadlock Avoidance

- ❖ Requires that the system has some additional ***a priori*** information available
- ❖ Simplest and most useful model requires that each process declare the ***maximum number*** of resources of each type that it may need
- ❖ The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- ❖ Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes

Safe State

- ❖ When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- ❖ System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the systems such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
- ❖ That is:
 - ❖ If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - ❖ When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - ❖ When P_i terminates, P_{i+1} can obtain its needed resources, and so on



johnjose@iitg.ac.in
<http://www.iitg.ac.in/johnjose/>